

# NODE.JS IN PRODUCTION





Dienstag, 31. August 2010









Dienstag, 31. August 2010







# COMMUNITY

- lomography.com (shameless plug)
- they upload their photos (~5000 per day)
- we show them on our site

# MANY DIFFERENT SIZES

- we need photos in all kind of sizes (thumbnail, grid, detail etc.)
- we don't know all needed sizes in advance
  - not every photo in every size
- we need to resize them on demand



# ON DEMAND RESIZING

- 1st attempt aka “lets start”
- sinatra/passenger application
  - queue resize job and process in background
- show “smaller version”
  - sometimes thumbnail







# ON DEMAND RESIZING

- 2nd attempt aka “the worst idea i’ve ever had”

- sinatra/passenger application

**40 instances**

- resize photo and deliver

**blocks normal photo  
delivery**

- blocks the entire server for



# ASYNCTO THE RESCUE

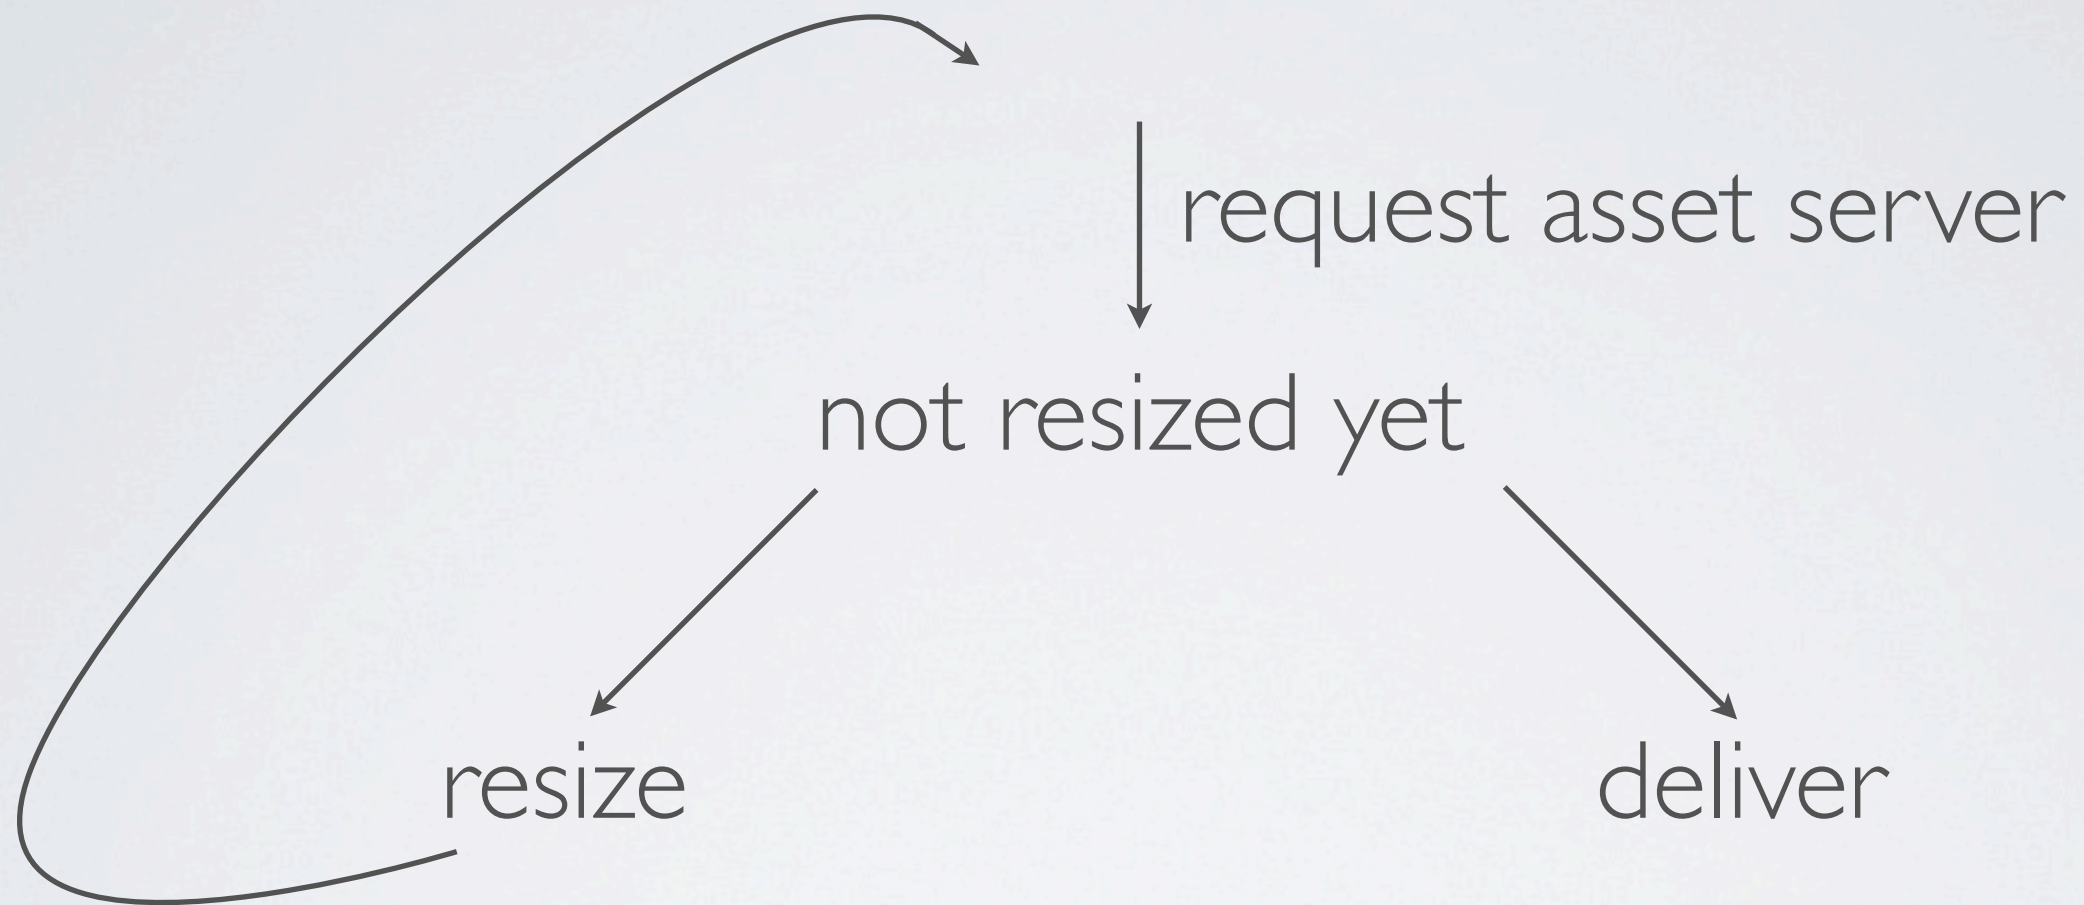
- i tried async sinatra with eventmachine
- can't read my own code after 6 hours
- node to the rescue
  - easier to write async code in javascript



# ON DEMAND RESIZING

- 3rd attempt aka “node rules them all”
  - node application
    - redirect to node server
    - resize photo
    - redirect back







# WHY NODE?

- a shitload of IO
  - check if file exists on disk (shared network device) (200 ms)
  - check if file exists on amazon S3 (upto 1 sec)
    - download original from Amazon S3 (upto 2 sec)



# SLOW START

- installed node
- built first prototype (happy face)
- tried some test “suites”
  - unable to install a single one
- plain javascript file test.js
  - “node test.js”



# GETTING READY “TESTING”

- installed on linux (our servers run linux)
- crashed immediately (sad face)

m(



# WHATS WRONG?

```
process.addListener('uncaughtException', function (err) {  
    console.log('Caught exception: ' + sys.inspect(err));  
    response.simpleText(500, "Internal Error");  
});
```

# READ THE CHANGELOG

- mac homebrew version 0.1.100
- debian package version 0.1.99
- “Introduce console.log() and friends.”

m(



# CALCULATE SHA1

```
var calculateSha1ForFile = function(filename, callback) {  
  var readStream = fs.createReadStream(filename);  
  var sha1ForFile = crypto.createHash('sha1');  
  
  readStream.addListener('error', function(err) {  
    callback(err);  
  });  
  readStream.addListener('data', function(data) {  
    sha1ForFile.update(data);  
  });  
  readStream.addListener('close', function() {  
    callback(false, sha1ForFile.digest('hex'));  
  });  
};
```

# CONVERT BINARY TO BINARY

```
readStream.addListener('data', function(data) {  
    sha1ForFile.update(data);  
});
```

```
readStream.addListener('data', function(data) {  
    sha1ForFile.update(data.toString('binary'));  
});
```

m(



# TOO MANY OPEN FILES

- http request library does not close connections (on linux)
- node IPv4 TCP **unknown protocol** (ESTABLISHED)

# REVERSE PROXY

- no security analysis (until now)
- do not connect node directly to the net
- use nginx or lighttpd (or some other evented webserver)
  - do some filtering



# RELEASE CYCLE

- new version every 1-2 weeks
- upgrade is a must
  - a lot of bugfixes
- you need tests! really!

# MONITORING

- node crashes sometimes (i don't know why)
- install some kind of monitoring (god or monit)
  - restart node after crash



# WAS IT WORTH?

- YES! node rocks!
- hard to write synchronous code
- resize photos and delivery normal thumbnails without 10000 instances aka “less memory”
- able to resize 5 photos per second
  - cpu utilization ~90%
- only 120 lines of javascript

@beanieboi



- <http://www.lomography.com/photos/6035921>
- <http://www.lomography.com/photos/11460174>
- <http://www.lomography.com/photos/6177726>